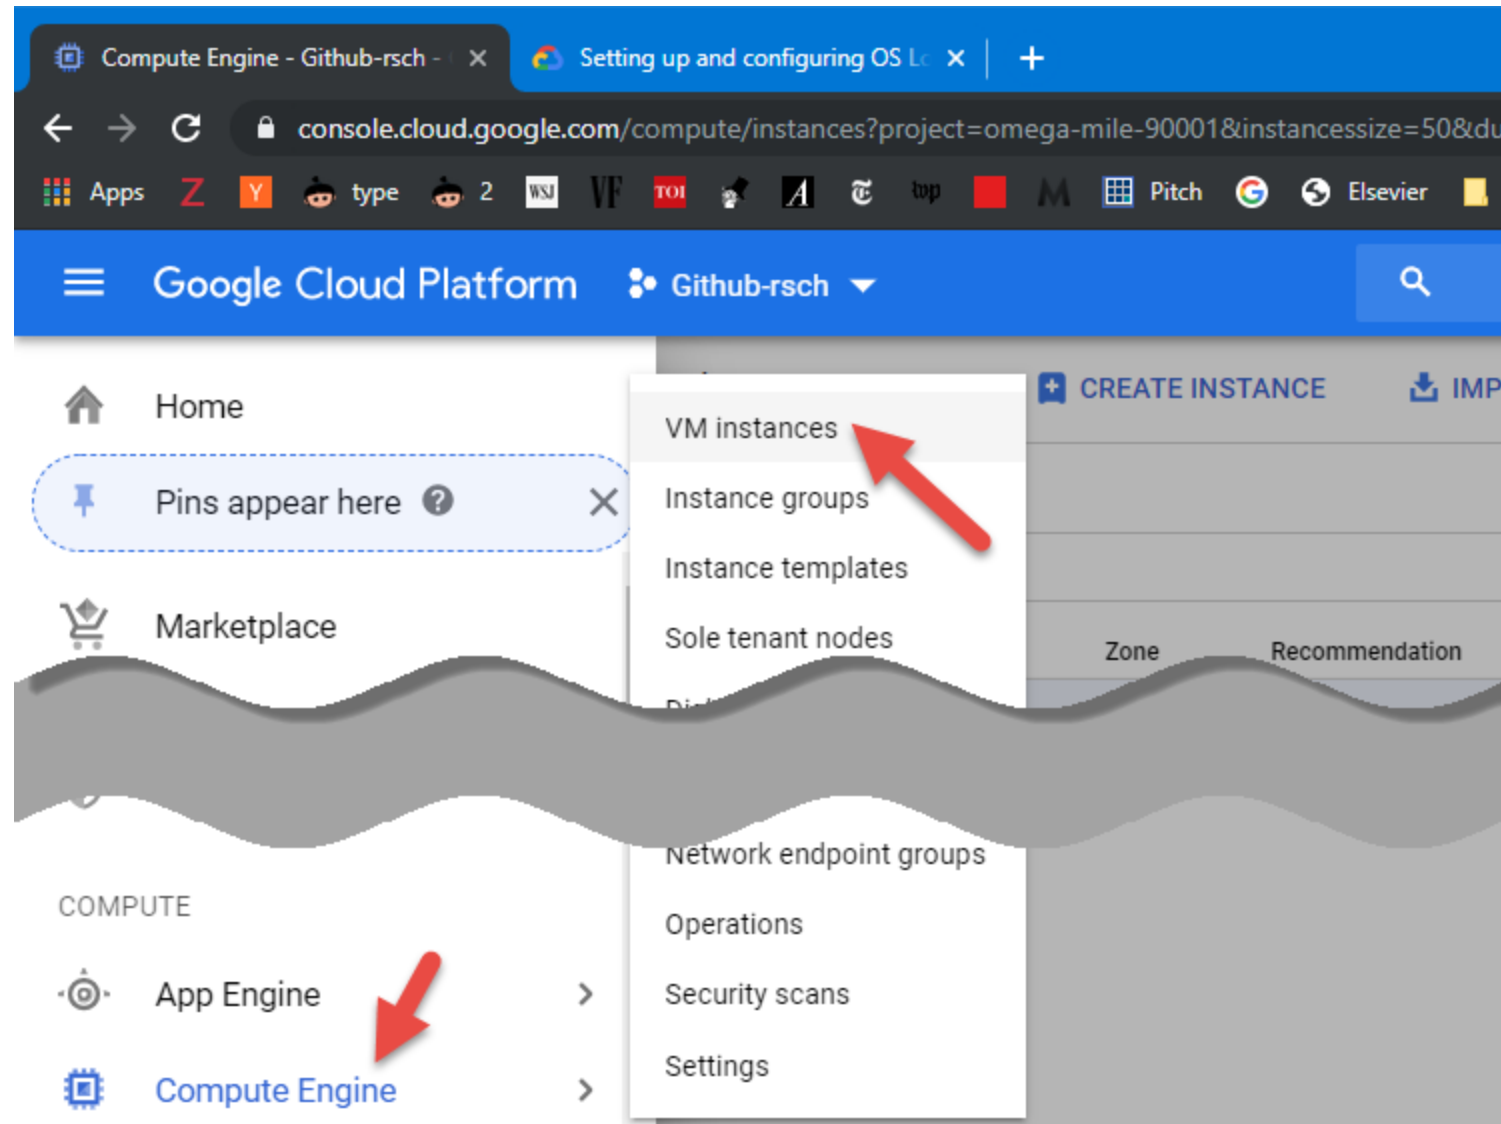


Google Cloud

Ramesh Shankar

University of Connecticut



Google Cloud Platform

Github-rsch

Compute Engine

VM instances

Instance groups

Instance templates

Sole tenant nodes

Disks

Snapshots

Images

VM instances

+

CREATE INSTANCE

IMPORT VM

REFRESH

Filter VM instances

<input type="checkbox"/>	Name ^	Zone	Recommendation	In use by	Internal IP
<input type="checkbox"/>	<input checked="" type="checkbox"/> instance-1-big-data-nyu-msba	us-east1-b			10.142.0.2 (nic0)
<input type="checkbox"/>	<input checked="" type="checkbox"/> instance-2-yelp	us-east4-a			10.150.0.3 (nic0)
<input type="checkbox"/>	<input checked="" type="checkbox"/> ramesh-regression-1	us-east1-b			10.142.0.2 (nic0)
<input type="checkbox"/>	<input checked="" type="checkbox"/> ras1-github-instance-1	us-east4-c			10.150.0.2 (nic0)



Create an instance

To create a VM instance, select one of the options:



New VM instance

Create a single VM instance from scratch



New VM instance from template

Create a single VM instance from an existing template



Marketplace

Deploy a ready-to-go solution onto a VM instance



You have a draft that wasn't submitted, click Restore to keep working on it

Restore

Name ?

instance-1

Region ?

us-central1 (Iowa)

Zone ?

us-central1-a

Machine configuration

Machine family

General-purpose

Memory-optimized

Machine types for common workloads, optimized for cost and flexibility

Generation

First

Powered by Skylake CPU platform or one of its predecessors

Machine type

n1-standard-1 (1 vCPU, 3.75 GB memory)



vCPU

1

Memory

3.75 GB

Container ?

☐ Deploy a container image to this VM instance. [Learn more](#)

Boot disk ?



New 10 GB standard persistent disk

Image

Debian GNU/Linux 9 (stretch)

Change

Identity and API access ?

Service account ?

Compute Engine default service account

Access scopes ?

- ☐ Allow default access
- ☒ Allow full access to all Cloud APIs
- ☐ Set access for each API

Firewall ?

Add tags and firewall rules to allow specific network traffic from the Internet

- ☐ Allow HTTP traffic
- ☐ Allow HTTPS traffic

[Management, security, disks, networking, sole tenancy](#)

You will be billed for this instance. [Compute Engine pricing](#) [↗](#)

Create

Cancel

Equivalent [REST](#) or [command line](#)

Google Cloud Platform Github-rsch

Compute Engine

VM instances

Filter VM instances

Name	Zone	Connect
<input type="checkbox"/> instance-1-big-data-nyu-msba	us-east1-b	SSH
<input type="checkbox"/> instance-2-yelp	us-east4-a	SSH
<input type="checkbox"/> ramesh-regression-1	us-east1-b	SSH
<input type="checkbox"/> ras1-github-instance-1	us-east4-c	SSH

1. Note the instance name
2. Note the zone
3. Click on the SSH button to start your instance

Cloud Platform

Github-rsch

Engine

VM instances

Filter VM instances

Name

☒ instance-1-big-data-nyu-msba

☐ instance-2-ye

☐ ramesh-regre

☐ ras1-github-i

discounts

ramesh_uconn@instance-1-big-data-nyu-msba: ~

* Documentation: <https://help.ubuntu.com>

* Management: <https://landscape.canonical.com>

* Support: <https://ubuntu.com/advantage>

System information as of Mon Sep 9 18:42:53 UTC 2019

System load: 0.11

Usage of /: 19.8% of 96.75GB

Memory usage: 1%

Swap usage: 0%

Processes: 121

Users logged in: 0

IP address for ens4: 10.142.0.2

* Congrats to the Kubernetes community on 1.16 beta 1! Now available in MicroK8s for evaluation and testing, with upgrades to RC and GA

snap info microk8s

30 packages can be updated.

0 updates are security updates.

Last login: Mon Sep 9 18:24:18 2019 from 74.125.177.165

ramesh_uconn@instance-1-big-data-nyu-msba:~\$

```
ramesh_uconn@instance-1-big-data-nyu-msba:~/taxi$ more yellowtaxi_2017.sh
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-01.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-02.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-03.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-04.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-05.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-06.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-07.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-08.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-09.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-10.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-11.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-12.csv
ramesh_uconn@instance-1-big-data-nyu-msba:~/taxi$ █
```



```
ramesh_uconn@instance-1-big-data-nyu-msba:~/taxi$ more yellowtaxi_2018.sh
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-01.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-02.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-03.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-04.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-05.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-06.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-07.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-08.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-09.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-10.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-11.csv
wget https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-12.csv
```

ramesh_uconn@instance-1-big-data-nyu-msba:~/taxi\$ sh yellowtaxi_2017.sh

1

```
--2019-09-09 18:55:15-- https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-01.csv
Resolving s3.amazonaws.com (s3.amazonaws.com)... 52.216.97.181
Connecting to s3.amazonaws.com (s3.amazonaws.com)|52.216.97.181|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 854903002 (815M) [application/x-www-form-urlencoded]
Saving to: 'yellow_tripdata_2017-01.csv'
```

yellow_tripdata_2017-01.c 100%[=====>] 815.30M 48.9MB/s in 15s

2

2019-09-09 18:55:31 (54.1 MB/s) - 'yellow_tripdata_2017-01.csv' saved [854903002/854903002]

```
--2019-09-09 18:55:31-- https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2017-02.csv
Resolving s3.amazonaws.com (s3.amazonaws.com)... 52.216.176.141
Connecting to s3.amazonaws.com (s3.amazonaws.com)|52.216.176.141|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 808065449 (771M) [application/x-www-form-urlencoded]
Saving to: 'yellow_tripdata_2017-02.csv'
```

yellow_tripdata_2017-02.c 21%[=====>] 163.48M 59.7MB/s

3

```
ramesh_uconn@instance-1-big-data-nyu-msba:~/taxi$ sh yellowtaxi_2018.sh
```

1

```
--2019-09-09 18:58:29-- https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-01.csv
Resolving s3.amazonaws.com (s3.amazonaws.com)... 52.216.105.141
Connecting to s3.amazonaws.com (s3.amazonaws.com)|52.216.105.141|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 772098307 (736M) [text/csv]
Saving to: 'yellow_tripdata_2018-01.csv'
```

2

```
yellow_tripdata_2018-01.c 100%[=====>] 736.33M 69.3MB/s in 8.9s
```

```
2019-09-09 18:58:38 (82.9 MB/s) - 'yellow_tripdata_2018-01.csv' saved [772098307/772098307]
```

```
--2019-09-09 18:58:38-- https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2018-02.csv
Resolving s3.amazonaws.com (s3.amazonaws.com)... 52.216.105.141
Connecting to s3.amazonaws.com (s3.amazonaws.com)|52.216.105.141|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 748827487 (714M) [text/csv]
Saving to: 'yellow_tripdata_2018-02.csv'
```

3

```
yellow_tripdata_2018-02.c 29%[=====>] 207.99M 90.4MB/s
```

Google Cloud Platform

Github-rsch

Home

Pins appear here ?

Kubernetes Engine

Cloud Functions

Cloud Run

STORAGE

Bigtable

Datastore

Firestore

Filestore

Storage

SQL

Spanner

Memorystore

NETWORKING

M instances

Filter VM instances

Name

instance-1-big-data-nyu-msba

instance-2-yelp

ramesh-regression-1

ras1-github-instance-1

Browser

Transfer

Transfer Appliance

Settings



- **Name your bucket**

Pick a globally unique, permanent name. [Naming guidelines](#)

1

Tip: Don't include any sensitive information

[CONTINUE](#)

- **Choose where to store your data**
- **Choose a default storage class for your data**
- **Choose how to control access to objects**
- **Advanced settings (optional)**

2

[CREATE](#)[CANCEL](#)

ramesh_uconn@instance-1-big-data-nyu-msba:~/taxi\$ more yellowtaxi_2017_2018_upload.sh

```
gsutil cp yellow_tripdata_2017-01.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2017-02.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2017-03.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2017-04.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2017-05.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2017-06.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2017-07.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2017-08.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2017-09.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2017-10.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2017-11.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2017-12.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2018-01.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2018-02.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2018-03.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2018-04.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2018-05.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2018-06.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2018-07.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2018-08.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2018-09.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2018-10.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2018-11.csv us-east1 gs://rs-nyc-taxi-1
gsutil cp yellow_tripdata_2018-12.csv us-east1 gs://rs-nyc-taxi-1
```

```
ramesh_uconn@instance-1-big-data-nyu-msba:~/taxi$ sh yellowtaxi_2017_2018_upload.sh
```

```
Copying file:///yellow_tripdata_2017-01.csv [Content-Type=text/csv]...  
==> NOTE: You are uploading one or more large file(s), which would run  
significantly faster if you enable parallel composite uploads. This  
feature can be enabled by editing the  
"parallel_composite_upload_threshold" value in your .boto  
configuration file. However, note that if you do this large files will  
be uploaded as `composite objects  
<https://cloud.google.com/storage/docs/composite-objects>`, which  
means that any user who downloads such objects will need to have a  
compiled crcmod installed (see "gsutil help crcmod"). This is because  
without a compiled crcmod, computing checksums on composite objects is  
so slow that gsutil disables downloads of composite objects.
```

```
CommandException: No URLs matched: us-east1/s
```

```
Copying file:///yellow_tripdata_2017-02.csv [Content-Type=text/csv]...  
==> NOTE: You are uploading one or more large file(s), which would run  
significantly faster if you enable parallel composite uploads. This
```

Google Cloud Platform

Github-rsch

20

Storage

Browser

Transfer

Transfer Appliance

Settings

← Bucket details

EDIT BUCKET

REFRESH BUCKET







rs-nyc-taxi-1

ObjectsOverviewPermissionsBucket Lock

Upload filesUpload folderCreate folderManage holdsDelete

Filter by prefix...

Buckets / rs-nyc-taxi-1

<input type="checkbox"/>	Name	Size	Type	Storage class	Last modified	Public access	Encryption	Retention expiration date	Holds
<input type="checkbox"/>	 yellow_tripdata_2017-01.csv	815.3 MB	text/csv	Standard	9/9/19, 3:13:27 PM UTC-4	Not public	Google-managed key	–	None
<input type="checkbox"/>	 yellow_tripdata_2017-02.csv	770.63 MB	text/csv	Standard	9/9/19, 3:13:45 PM UTC-4	Not public	Google-managed key	–	None
<input type="checkbox"/>	 yellow_tripdata_2017-03.csv	865.56 MB	text/csv	Standard	9/9/19, 3:14:04 PM UTC-4	Not public	Google-managed key	–	None
<input type="checkbox"/>	 yellow_tripdata_2017-04.csv	844.61 MB	text/csv	Standard	9/9/19, 3:14:20 PM UTC-4	Not public	Google-managed key	–	None
<input type="checkbox"/>	 yellow_tripdata_2017-05.csv	849.68 MB	text/csv	Standard	9/9/19, 3:14:40 PM UTC-4	Not public	Google-managed key	–	None
<input type="checkbox"/>	 yellow_tripdata_2017-06.csv	812.44 MB	text/csv	Standard	9/9/19, 3:14:59 PM UTC-4	Not public	Google-managed key	–	None



Google Cloud Platform



Github-rsch



Home



Pins appear here ?



Composer



Dataproc



Pub/Sub



Dataflow



BigQuery



IoT Core



Data Catalog



Data Fusion



Genomics



Healthcare



SHORTCUTS

Unsaved query

1



Run



omega-mile-90



BigQuery

FEATURES & INFO

SHORTCUTS

[+ COMPOSE NEW QUERY](#)

Query history

Saved queries

Job history

Transfers

Scheduled queries

BI Engine

Resources

[+ ADD DATA](#)

▼ omega-mile-90001

▶ greentaxi_msbapm1

yelp_business

▶ alien-climber-851

▶ bigquery-public-data

▶ bigquery-samples

▶ fh-bigquery

Unsaved query Edited

[HIDE EDITOR](#)[FULL SCREEN](#)

1



Run



Save query



omega-mile-90001

[+ CREATE DATASET](#)[UNPIN PROJECT](#)

Create dataset

Dataset ID

1

nyctaxi1

Data location (Optional) ?

2

United States (US)

Default table expiration ?

☒ Never

☐ Number of days after table creation:

Encryption

Data is encrypted automatically. Select an encryption key management solution.

☒ Google-managed key

No configuration required

☐ Customer-managed key

Manage via Google Cloud Key Management Service

3

Create dataset

Cancel

Google Cloud Platform

Github-rsch

20

BigQuery

FEATURES & INFO

SHORTCUTS

+ COMPOSE NEW QUERY

Query history

Saved queries

Resources

+ ADD DATA

Search for your tables and datasets

omega-mile-90001

greentaxi_msbapm1

nyctaxi1

yelp_business

Unsaved query Edited

HIDE EDITOR

FULL SCREEN

1

Run Save query

2

omega-mile-90001:nyctaxi1

CREATE TABLE

SHARE DATASET

COPY DATASET

DELETE DATASET

Create table

Source

Create table from:

Google Cloud Storage

Select file from GCS bucket: ?

✓ rs-nyc-taxi-1/*

Browse

File format:

CSV

Destination

Project name

Github-rsch

Dataset name

nyctaxi1

Table type ?

Native table

Table name

nyc_taxi_data_1

Schema

Auto detect

☒ Schema and input parameters

i Schema will be automatically generated.

Partition and cluster settings

Partitioning: ?

No partitioning

Advanced options

Create table

Cancel



Query history

Saved queries

Resources

[+ ADD DATA](#) ▼

▼ omega-mile-90001

▶ greentaxi_msbapm1

▼ nyctaxi1

nyc_taxi_data_1

yelp_business

▶ alien-climber-851

▶ bigquery-public-data

▶ bigquery-samples

Query editor

1 `SELECT count(*) FROM `omega-mile-90001.nyctaxi1.nyc_taxi_data_1``

Processing location: US

Run ▼

Save query

Save view

Schedule query ▼

More ▼

Query results

[SAVE RESULTS](#) ▼ [EXPLORE WITH DATA STUDIO](#)

Query complete (0.9 sec elapsed, 0 B processed)

Job information

Results

JSON

Execution details

Row f0_

1 216301148

```
ramesh_uconn@instance-1-big-data-nyu-msba:~/taxi$ wc -l yellow_tripdata_2018-11.csv
8145166 yellow_tripdata_2018-11.csv
ramesh_uconn@instance-1-big-data-nyu-msba:~/taxi$ wc -l *.csv
 9710126 yellow_tripdata_2017-01.csv
 9168827 yellow_tripdata_2017-02.csv
10294630 yellow_tripdata_2017-03.csv
10046190 yellow_tripdata_2017-04.csv
10102126 yellow_tripdata_2017-05.csv
 9656995 yellow_tripdata_2017-06.csv
 8588488 yellow_tripdata_2017-07.csv
 8422155 yellow_tripdata_2017-08.csv
 8945461 yellow_tripdata_2017-09.csv
 9768817 yellow_tripdata_2017-10.csv
 9284805 yellow_tripdata_2017-11.csv
 9508278 yellow_tripdata_2017-12.csv
 8759876 yellow_tripdata_2018-01.csv
 8492078 yellow_tripdata_2018-02.csv
 9430378 yellow_tripdata_2018-03.csv
 9305517 yellow_tripdata_2018-04.csv
 9224065 yellow_tripdata_2018-05.csv
 8713833 yellow_tripdata_2018-06.csv
 7849750 yellow_tripdata_2018-07.csv
 7849136 yellow_tripdata_2018-08.csv
 8040135 yellow_tripdata_2018-09.csv
 8821107 yellow_tripdata_2018-10.csv
 8145166 yellow_tripdata_2018-11.csv
 8173233 yellow_tripdata_2018-12.csv
216301172 total
```



Query history

Saved queries

Job history

Resources

+ ADD DATA

Search for your tables and datasets

omega-mile-90001

greentaxi_msbapm1

nyctaxi1

nyc_taxi_data_1

yelp_business

alien-climber-851

bigquery-public-data

bigquery-samples

fh-bigquery

githubarchive

imjasonh-storage

Unsaved query Edited

1

```
1 SELECT VendorID, sum(fare_amount) as total_fare
2 FROM `omega-mile-90001.nyctaxi1.nyc_taxi_data_1`
3 group by VendorID
```

Processing location: US

Run

Save query

Save view

Schedule query

More

Query results

SAVE RESULTS

EXPLORE WITH DATA STUDIO

Query complete (1.5 sec elapsed, 3.2 GB processed)

3

Job information

Results

JSON

Execution details

Row	VendorID	total_fare
1	1	1.2102583026599963E9
2	2	1.6070634624099953E9
3	4	6210471.119999999
4	null	null

2